

Authors' Responses

Journal Title: IEEE Robotics and Automation Letters
Reference Number: 26-3841v1
Title of Paper: VLM-DexSafeDagger: VLM-Guided Intervention Policy Towards Safer Visuomotor Skill Distillation
Authors: Anonymous Authors

We sincerely thank the Editor and the reviewers for their careful evaluation of our manuscript and for their constructive comments. Their feedback has helped us identify where the contribution, safety claims, experimental evidence, and implementation details require clearer presentation.

We are preparing a major revision of the manuscript. In the revised version, we will:

- 1) clarify the specific contribution of VLM-guided threshold adaptation through comparisons with conventional fixed and adaptive alternatives;
- 2) revise the safety terminology and explicitly delimit the empirical, task-specific scope of the reported safety results;
- 3) strengthen the evaluation of the Bellman critic, teacher policies, real-world experiments, and resource usage; and
- 4) provide the implementation details needed to reproduce the VLM-based threshold tuner and the intervention pipeline.

Detailed responses to each comment are provided below. Reviewer comments are reproduced in black, and our responses are shown in blue. Red bracketed text marks details that will be completed after the corresponding revision or experiment has been finalized. Changes in the revised manuscript will be highlighted in blue.

1 Response to Reviewer #14

We sincerely thank the reviewer for the careful assessment of our work and for the constructive suggestions. Below are our point-by-point responses.

Comment 1: *whether the VLM and visual information provide a clear advantage over simpler adaptive approaches.*

[Response:] We have added ablations (g) and (h) to isolate the contribution of the VLM to the adaptive intervention pipeline. We have also expanded the discussion of Table III in the Results section. The results provide strong evidence that VLM-based intervention outperforms both LLM-based intervention and performance-based curriculum adaptation. Semantic information alone reduces the unsafe exploration rate, while combining visual and semantic information enables more precise intervention, achieving the highest task success rate and the lowest unsafe exploration rate. Although the performance-based curriculum reduces unsafe exploration to some extent, it sacrifices task success, suggesting that a simple rule-based method is insufficient for guiding intervention policies in complex scenarios.

Comment 2: *The related work on adaptive intervention should also be expanded.*

[Response:] We have reformulated the Related Work section to more clearly discuss adaptive intervention methods. In particular, AIM provides an important foundation for our work, and we have added it to the reference list. Compared with prior methods such as HG-Dagger and ThriftyDagger, our proposed pipeline does not require human intervention, enabling fully autonomous intervention during online policy distillation.

Comment 3: *the authors should provide the VLM prompt and explain how the images are selected and how the VLM outputs are used.*

[Response:] The details required to reproduce our method are provided in the supplementary material. The exact VLM prompt and representative visual input examples are included in Supplementary Section B.

Comment 4: *The Bellman success-value critic is also important for intervention, but its performance is not evaluated in detail. Some evaluation of the critic, especially its errors for unsafe cases, would be helpful.*

[Response:] We evaluate the resulting student and fixed-threshold Bellman critic from ablation (e) and report the recorded metrics in Supplementary Section C.1. Both networks are frozen, and teacher intervention is disabled so that the student remains in control throughout each evaluation episode. The critic runs in parallel

as a passive evaluator whose predictions do not affect the executed actions. This setup preserves observed unsafe terminations and allows us to evaluate whether the critic signals low predicted safe-lift success before an unsafe event occurs.

Comment 5: *Please also clarify the number of runs/random seeds used for the results.*

[Response:] We run 30 real-world evaluations for each object, distributed equally across random seeds 40, 41, 42, 43, and 44 (six evaluations per seed).

Comment 6: *Finally, I am not fully convinced by the use of the term “safety-critical.” The current safety evaluation is based on task-specific unsafe conditions rather than formal safety guarantees. The authors should clarify the scope of this claim.*

[Response:] We acknowledge that the term “safety-critical” is too strong for an adaptive intervention pipeline such as ours. We have therefore revised the title to “VLM-DexSafeDagger: VLM-Guided Intervention Policy Towards Safer Visuomotor Skill Distillation” and removed or qualified the term elsewhere in the manuscript. The revised text explicitly states that our method provides empirical risk reduction under task-specific event definitions rather than a formal safety guarantee, and that safety violations remain possible.

Comment 7: *However, only VLM-DexSafeDagger evaluated on the real robot does not sufficiently prove the proposed concept.*

[Response:] We deploy the frozen RGB student policies resulting from ablations (e) and (f) on the real robot using the same object set, initial-condition protocol, and number of trials. We report lift-success and unsafe-event rates for the seen and unseen object sets in Supplementary Section C.4. Only the resulting student policy is deployed; the teacher, Bellman critic, VLM, and threshold-adaptation mechanism are not active during physical evaluation.

Comment 8: *It would be helpful to provide overall results for all 22 objects, for example by reporting results separately for seen and unseen objects.*

[Response:] We have doubled the number of evaluation episodes for each object and reported the per-object and aggregate seen/unseen results for all 22 objects in Supplementary Tables V and VI. As recommended, we have also extended Tables II–IV in the main paper to report aggregate results for each evaluation setting.

2 Response to Reviewer #22

We sincerely thank the reviewer for the detailed and constructive comments. Below are our point-by-point responses.

Comment 1: *The term “safety-critical” is too strong because the framework supplies empirical intervention rather than formal constraint satisfaction. State explicitly that safety violations remain possible.*

[Response:] We acknowledge that the term “safety-critical” is too strong for an adaptive intervention pipeline such as ours. We have therefore revised the title to “VLM-DexSafeDagger: VLM-Guided Intervention Policy Towards Safer Visuomotor Skill Distillation” and removed or qualified the term elsewhere in the manuscript. The revised text explicitly states that our method provides empirical risk reduction under task-specific event definitions rather than a formal safety guarantee, and that safety violations remain possible.

Comment 2: *Demonstrate that the VLM is necessary by comparing against fixed thresholds, rule-based adaptation, Bayesian optimization, curriculum schedules, grid search, and a small learned threshold network.*

[Response:] We have supplemented the ablation study with variants (g) and (h). Together with variants (e) and (f), these experiments form a controlled ladder of adaptive intervention strategies: fixed thresholds, a performance-based curriculum, semantics-guided intervention-policy adaptation, and visual-semantic-guided adaptive intervention. We have added a corresponding discussion to the Results section. Table III shows that adaptive intervention reduces the unsafe rate relative to fixed thresholds. However, the curriculum-based method sacrifices lift success, whereas incorporating semantic information recovers task performance while further reducing unsafe exploration. Overall, combining visual and semantic information outperforms the other approaches, achieving the best balance between lift success and unsafe rate.

Comment 3: *The Bellman critic predicts combined “safe-lift-success,” thereby conflating task success with safety. Separate safety risk and task value, and report calibration, false-negative rates, and reliability diagrams.*

[Response:] Lift success and unsafe termination measure different aspects of an episode and are not mutually exclusive. Lift success is recorded once the object satisfies the prescribed height-and-hold criterion, but satisfying this criterion does not terminate the rollout. The episode continues to be monitored for unsafe events until timeout or unsafe termination. Consequently, an episode may achieve lift success and subsequently terminate because of an unsafe event. The Bellman critic is therefore trained on the stricter joint safe-lift-success target, which is positive only when the lift criterion is satisfied without an unsafe event.

We evaluate the resulting student and fixed-threshold Bellman critic from ablation (e). Both networks are frozen, teacher intervention is disabled, and the critic runs in parallel as a passive evaluator. We separately report task success and unsafe termination outcomes, including the critic’s unsafe-episode detection and false-negative rates. The results are reported in Supplementary Section C.1.

Comment 4: *Teacher intervention is assumed safe, yet teacher policies may collide or fail. Quantify teacher unsafe rates and define how the framework responds when teacher and critic disagree.*

[Response:] Teacher takeover is not guaranteed to be safe; it provides an empirical tendency toward safer exploration rather than a formal guarantee. We report the teacher policies’ unsafe rates in Table III of the main paper and, in this revision, provide mean statistics across all tested objects. The intervention gate follows a conservative OR rule: intervention is triggered if either the teacher–student action disagreement exceeds its threshold or the critic’s predicted success falls below its threshold. Thus, when the two signals disagree, the conservative outcome is to intervene. Comparing ablations (c) and (e) shows that adding the critic-based branch substantially reduces the unsafe exploration rate.

Comment 5: *Provide exact VLM prompts, model version, decoding settings, image-selection procedure, output parsing, update frequency, latency, cost, and failure handling. Reproducibility may otherwise depend on an external proprietary service.*

[Response:] The details required to reproduce our method are provided in the supplementary material. The exact VLM prompt, representative visual-input examples, model and inference configuration, image-selection and output-processing procedures, update frequency, token usage, and failure handling are included in Supplementary Section B.

Comment 6: *Real-world unseen-object performance is weak: lift success falls to 60–66.7%, while unsafe episodes reach 40%. Claims of sufficient deployment robustness should be moderated.*

[Response:] We will remove the claim that the current results establish sufficient robustness for deployment.

Comment 7: *Only fifteen real-world episodes per reported object produce coarse estimates. Increase repetitions and report confidence intervals and statistical comparisons across all 22 objects.*

[Response:] We have doubled the number of real-world evaluation episodes from 15 to 30 per object. We have also reported the complete results for all 22 tested objects in Supplementary Section C.3. Supplementary Table VI reports the pooled means and Wilson 95% confidence intervals for the 14 seen and eight unseen objects.

Comment 8: *Safety-event thresholds for harmful collision, object escape, palm flipping, and workspace*

violation require physical justification. Include contact-force measurements and severity-weighted safety metrics.

[Response:] We have provided the exact definitions of all unsafe events in the termination-conditions discussion in Supplementary Section A.2. When multiple unsafe conditions occur simultaneously during training, the event label is assigned according to a clear severity-based priority: harmful collision, object escape, palm flipped, and hand too far, in descending order.

Comment 9: *Clarify whether VLM adaptation occurs only during simulation distillation or during physical deployment. If absent at deployment, the role of scene-level VLM reasoning should not be overstated.*

[Response:] The VLM is queried only during simulation-based online distillation. During physical deployment, the RGB student acts without the VLM, privileged teacher, Bellman critic, or threshold adaptation. We will state this distinction in the abstract, method, experimental setup, and limitations, and will describe the VLM as a training-time adaptation mechanism rather than a deployment-time safety monitor.

Comment 10: *Compare training samples, teacher queries, interventions, computation, and final success across Vanilla DAgger, SafeDAgger, critic-based intervention, and VLM-guided intervention using multiple random seeds.*

[Response:] We repeat the complete distillation process for Vanilla DAgger (b), Vanilla SafeDAgger (c), critic-based intervention (e), and VLM-guided intervention (f) using five independent random seeds, {40, 41, 42, 43, 44}. All variants use the same number of training samples, and the teacher-intervention rates exhibit only minor variation across seeds. The final lift-success and unsafe-episode rates are reported as the mean $\pm$ standard deviation across seeds in Supplementary Section C.2.

We thank the Editor and the reviewers again for their thoughtful feedback. The requested revisions will help us present the method’s contribution more clearly, support its empirical claims more rigorously, and state its limitations more precisely.